

CB19342 – COMPUTATIONAL STATISTICS LAB MANUAL



# **RAJALAKSHMI ENGINEERING COLLEGE**

**An AUTONOMOUS Institution  
Affiliated to ANNA UNIVERSITY, Chennai**

## **DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING**

### **CB19342 – COMPUTATIONAL STATISTICS LAB MANUAL THIRD YEAR**

**FIFTH SEMESTER 2024 - 2020**

**ODD SEMESTER**

## **1. PREDICTING HOUSE PRICES**

<b>EX.N0 : 1</b>	<b>Predicting House Prices</b>
<b><u>DATE : 24/07/2024</u></b>	

**PROBLEM STATEMENT:** Build a regression model to predict house prices based on features like location, size, and amenities.

**PYTHON CONCEPTS:** Functions, classes, numeric types, sequences.

**VISUALIZATION:** Plotting regression line, residual plots.

**MULTIVARIATE ANALYSIS:** Multiple regression.

**DATASET:** Kaggle House Prices

### **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.


### **PROGRAM:**

```
import pandas as pd from sklearn.preprocessing
import LabelEncoder from sklearn.model_selection
import train_test_split from sklearn.linear_model
import LinearRegression from sklearn.metrics
```

```

import r2_score, mean_absolute_error import
matplotlib.pyplot as plt file_path =
'C:/Users/APPU/Downloads/Housing.csv'
housing_data = pd.read_csv(file_path)
categorical_features = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
'prefarea', 'furnishingstatus']
le = LabelEncoder() for feature in
categorical_features:
housing_data[feature] = le.fit_transform(housing_data[feature])
X = housing_data.drop('price', axis=1) y = housing_data['price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression() model.fit(X_train, y_train) y_pred =
model.predict(X_test) r2 = r2_score(y_test, y_pred) mae =
mean_absolute_error(y_test, y_pred)
plt.figure(figsize=(10, 6)) plt.scatter(y_test,
y_pred, alpha=0.7, color='b')
plt.plot([y_test.min(), y_test.max()],
[y_test.min(), y_test.max()], 'k--', lw=2)
plt.xlabel('Actual Price') plt.ylabel('Predicted
Price') plt.title('Actual vs Predicted Prices')
plt.show() print(f'R-squared ( $R^2$ ): {r2}')
print(f'Mean Absolute Error (MAE): {mae}')

```



```
import numpy as np
test=np.array([ 7420,4,2,3,1,0,0,0,1,2,1,0]).reshape(-12,12)
model.predict(test)

array([8004072.41154001])
```

### **RESULT:**

Thus, the program for house price prediction is executed successfully.

## **2. CUSTOMER SEGMENTATION FOR AN E-COMMERCE COMPANY**

<b>EX.N0 : 2</b>	<b>Customer Segmentation for an E-commerce Company</b>
------------------	--

**DATE : 00/08/2024**

**PROBLEM STATEMENT:** Perform cluster analysis to segment customers based on purchasing behaviour.

**PYTHON CONCEPTS:** Data structures, file reading/writing.

**VISUALIZATION:** Cluster plots.

**MULTIVARIATE ANALYSIS:** Cluster analysis with k-means, hierarchical clustering.

**DATASET:** Online Retail Dataset **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

**PROGRAM:**

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

import seaborn as sns
import os

os.environ['OMP_NUM_THREADS'] = '1'


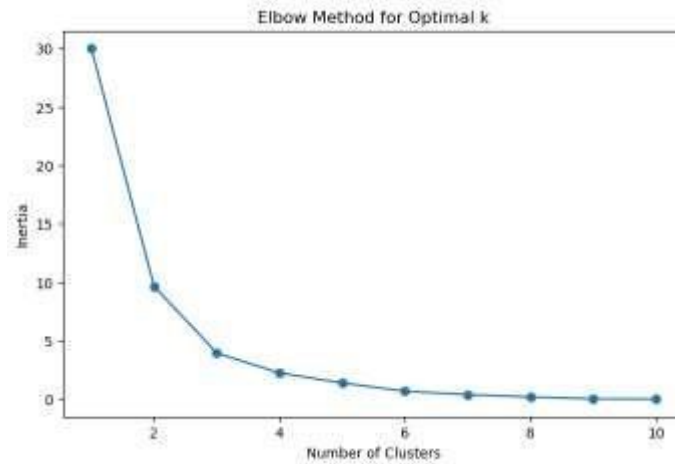
data = {'CustomerID': [1, 2, 3, 4, 0, 6, 7, 8, 9, 10],
        'Age': [20, 40, 30, 00, 23, 33, 43, 36, 29, 00],
        'AnnualIncome': [00000, 60000, 70000, 80000, 40000, 70000, 80000, 72000, 48000, 90000],
```

```

'SpendingScore': [60, 70, 80, 90, 00, 80, 90, 78, 60, 90] }
df = pd.DataFrame(data) features = df[['Age',
'AnnualIncome', 'SpendingScore']] scaler =
StandardScaler() scaled_features =
scaler.fit_transform(features) inertia = [] k_range =
range(1, 50) for k in k_range:
kmeans = KMeans(n_clusters=k, n_init=10, random_state=0) kmeans.fit(scaled_features)
inertia.append(kmeans.inertia_) plt.figure(figsize=(8, 0)) plt.plot(k_range, inertia, marker='o')
plt.xlabel('Number of Clusters') plt.ylabel('Inertia') plt.title('Elbow Method for Optimal k')
plt.show() optimal_k = 3 kmeans = KMeans(n_clusters=optimal_k, n_init=10,
random_state=0) df['Cluster'] = kmeans.fit_predict(scaled_features) plt.figure(figsize=(10,
7)) sns.scatterplot(data=df, x='AnnualIncome', y='SpendingScore', hue='Cluster',
palette='viridis', s=150) plt.title('Customer Segments') plt.xlabel('Annual Income')
plt.ylabel('Spending Score') plt.legend(title='Cluster') plt.show() print(df)

```

### **OUTPUT:**



### RESULT:

Thus, the program Segmentation for Company is successfully.

for Customer an E-commerce executed

### MOVIE REVIEWS

### 3. SENTIMENT ANALYSIS OF

EX.N0 : 3

## SENTIMENT ANALYSIS OF MOVIE REVIEWS

**DATE : 07/08/2024**

**PROBLEM STATEMENT:** Classify movie reviews as positive or negative using text Data.

**PYTHON CONCEPTS:** Text files, sequences, flow controls.

**VISUALIZATION:** Word cloud, bar plots.

**MULTIVARIATE ANALYSIS:** PCA for text data, logistic regression.

**DATASET:** IMDB Movie Reviews.

**ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 5: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

**PROGRAM:**

```
import pandas as pd
import matplotlib.pyplot as plt
from wordcloud import WordCloud
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer
import seaborn as sns
nltk.download('punkt')
nltk.download('stopwords')
df = pd.read_csv('C:/Users/AI_LAB/Downloads/IMDB Dataset.csv')
```



```

stop_words = set(stopwords.words('english')) stemmer =
PorterStemmer() def preprocess_text(text):
tokens = word_tokenize(text.lower()) tokens = [stemmer.stem(word) for word in tokens if
word.isalpha() and word not in stop_words] return ' '.join(tokens) df['cleaned_review'] =
df['review'].apply(preprocess_text) vectorizer = TfidfVectorizer(max_features=10000)
X = vectorizer.fit_transform(df['cleaned_review']).toarray() encoder =
LabelEncoder() y = encoder.fit_transform(df['sentiment']) pca =
PCA(n_components=2) X_pca = pca.fit_transform(X)
plt.figure(figsize=(8, 6)) plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y,
cmap='coolwarm', alpha=0.05) plt.title('PCA of Movie Reviews')
plt.xlabel('Principal Component 1') plt.ylabel('Principal Component
2') plt.colorbar(label='Sentiment') plt.show()
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) model
= LogisticRegression(max_iter=1500)

```



```


model.fit(X_train, y_train) y_pred = model.predict(X_test)
print("Confusion Matrix:") print(confusion_matrix(y_test, y_pred))
print("\nClassification Report:") print(classification_report(y_test,
y_pred)) positive_reviews = ''.join(df[df['sentiment'] ==
1]['cleaned_review']) negative_reviews = ''.join(df[df['sentiment']
== 0]['cleaned_review']) plt.figure(figsize=(12, 6)) if
len(positive_reviews.strip()) > 0:
plt.subplot(1, 2, 1) plt.imshow(WordCloud(width=800,
height=400,
background_color='white').generate(positive_reviews), interpolation='bilinear')
plt.title('Positive Reviews') plt.axis('off') else:
print("No content available for positive reviews.") if
len(negative_reviews.strip()) > 0:
plt.subplot(1, 2, 2) plt.imshow(WordCloud(width=800,
height=400,
background_color='white').generate(negative_reviews), interpolation='bilinear')
plt.title('Negative Reviews') plt.axis('off')
else:
print("No content available for negative
reviews.") plt.show() sns.countplot(x='sentiment',
data=df) plt.title('Sentiment Distribution')
plt.xlabel('Sentiment') plt.ylabel('Count')
plt.show()

```

---

**OUTPUT:**

COMPUTATIONAL STATISTICS



Confusion Matrix:

```
[[4306  655]
 [ 511 4528]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.87	0.88	4961
1	0.87	0.90	0.89	5039
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000

## **RESULT:**

Thus, the program for sentiment analysis of movie reviews is executed successfully.

#### **4. STOCK MARKET ANALYSIS**

<b>EX.N0 : 4</b>	<b>STOCK MARKET ANALYSIS</b>
<b><u>DATE : 14/08/2024</u></b>	

**PROBLEM STATEMENT:** Analyse stock market data to predict future stock prices.

**PYTHON CONCEPTS:** Data structures, file reading/writing, functions.

**VISUALIZATION:** Line plots, candlestick charts.

**MULTIVARIATE ANALYSIS:** Time series analysis, regression.

**DATASET:** Yahoo Finance Stock Data.

#### **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

#### **PROGRAM:**


```
import pandas as pd
import matplotlib.pyplot as plt
import mplfinance as mpf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
import numpy as np
file_path =
```

```

r'C:\Users\APPU\Downloads\yahoo_data.xlsx' data =
pd.read_excel(file_path, index_col='Date',
parse_dates=True) data.rename(columns={'Close*':
'Close', 'Adj Close**': 'Adj Close'}, inplace=True)
data.sort_index(inplace=True) data.ffill(inplace=True) if
'Adj Close' in data.columns:
plt.figure(figsize=(12, 6)) plt.plot(data['Adj Close'], label='Adjusted Close Price')
plt.title('Adjusted Close Price Over Time') plt.xlabel('Date') plt.ylabel('Price (USD)') plt.legend()
plt.show() reduced_data = data[-150:] # Reduce data points for candlestick chart
mpf.plot(reduced_data, type='candle', style='charles', title='Candlestick Chart') train_data,
test_data = data['Adj Close'][:int(len(data)*0.8)], data['Adj Close'][int(len(data)*0.8):] model =
ARIMA(train_data, order=(0, 1, 0)) model_fit = model.fit() forecast =
model_fit.forecast(steps=len(test_data)) mse = mean_squared_error(test_data, forecast) rmse =
np.sqrt(mse) print(f'RMSE: {rmse}') plt.figure(figsize=(12, 6))
plt.plot(train_data.index, train_data, label='Train Data')
plt.plot(test_data.index, test_data, label='Test Data')
plt.plot(test_data.index, forecast, label='Forecast')
plt.title('Stock Price Prediction') plt.xlabel('Date')
plt.ylabel('Price (USD)') plt.legend() plt.show()

```

## **OUTPUT:**



## **RESULT:**

Thus, the program for stock market analysis is executed successfully.

## **0. LOAN DEFAULT PREDICTION**



<b>EX.N0 : 0</b>	<b>LOAN DEFAULT PREDICTION</b>
<b><u>DATE : 21/08/2024</u></b>	

**PROBLEM STATEMENT:** Predict loan default probability based on borrower information.

**PYTHON CONCEPTS:** Classes, functions, sequences.

**VISUALIZATION:** ROC curve, bar plots.

**MULTIVARIATE ANALYSIS:** Logistic regression, factor analysis.

**DATASET:** Lending Club Loan Data **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

**PROGRAM:**


```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
import os
file_path = 'C:/Users/APPU/Downloads/loan_data.csv' # Update
```

```

path accordingly if os.path.exists(file_path): df =
pd.read_csv(file_path) print("Data loaded
successfully.") else: print(f"File not found:
{file_path}") dummies =
pd.get_dummies(df['purpose'], drop_first=True) df =
pd.concat([df, dummies], axis=1) df.drop('purpose',
inplace=True, axis=1) X = df.drop(['not.fully.paid'],
axis=1) y = df['not.fully.paid'] scaler =
StandardScaler() X_scaled = scaler.fit_transform(X)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size=0.33, random_state=42)
model = LogisticRegression() model.fit(X_train, y_train) y_pred_prob =
model.predict_proba(X_test)[:, 1] fpr, tpr, _ = roc_curve(y_test, y_pred_prob) roc_auc =
auc(fpr, tpr) plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--') plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.00]) plt.xlabel('False Positive Rate') plt.ylabel('True Positive
Rate')
plt.title('Receiver Operating Characteristic
(ROC)') plt.legend(loc='lower right') plt.show()

```

## **OUTPUT:**



### **RESULT:**

Thus, the program for loan default prediction is executed successfully.

## **6. IMAGE CLASSIFICATION**

<b>EX.N0 : 6</b>	<b>IMAGE CLASSIFICATION</b>
<b><u>DATE : 04/09/2024</u></b>	

**PROBLEM STATEMENT:** Classify images into categories using various features.

**PYTHON CONCEPTS:** File handling, classes.

**VISUALIZATION:** Image plots, feature importance plots.

**MULTIVARIATE ANALYSIS:** PCA, clustering.

**DATASET:** CIFAR-10 Dataset **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

### **PROGRAM:**

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt import numpy as np
```

```
(X_train, y_train), (X_test, y_test) = tf.keras.datasets.cifar10.load_data()
X_train, X_test = X_train / 200.0, X_test / 200.0 class_names =
['airplane', 'automobile', 'bird', 'cat', 'deer',
'dog', 'frog', 'horse', 'ship', 'truck']
```



```

plt.figure(figsize=(10,10)) for i in range(20): plt.subplot(0,0,i+1)
plt.xticks([]) plt.yticks([]) plt.grid(False) plt.imshow(X_train[i],
cmap=plt.cm.binary) plt.xlabel(class_names[y_train[i][0]]) plt.show()
model = models.Sequential([ layers.Conv2D(32, (3, 3),
activation='relu', input_shape=(32, 32, 3)), layers.MaxPooling2D((2,
2)), layers.Conv2D(64, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)), layers.Conv2D(64, (3, 3),
activation='relu'), layers.Flatten(), layers.Dense(64, activation='relu'),
layers.Dense(10) ]) model.compile(optimizer='adam',
loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
metrics=['accuracy']) history = model.fit(X_train, y_train, epochs=10,
validation_data=(X_test, y_test)) test_loss, test_acc =
model.evaluate(X_test, y_test, verbose=2) print(f"\nTest accuracy:
{test_acc}") plt.figure(figsize=(8, 4)) plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy']) plt.title('Model accuracy')
plt.ylabel('Accuracy') plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.subplot(1, 2, 2) plt.plot(history.history['loss'])
plt.plot(history.history['val_loss']) plt.title('Model
loss') plt.ylabel('Loss') plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.tight_layout() plt.show() predictions =
model.predict(X_test)
plt.figure(figsize=(10, 10))

```

COMPUTATIONAL STATISTICS

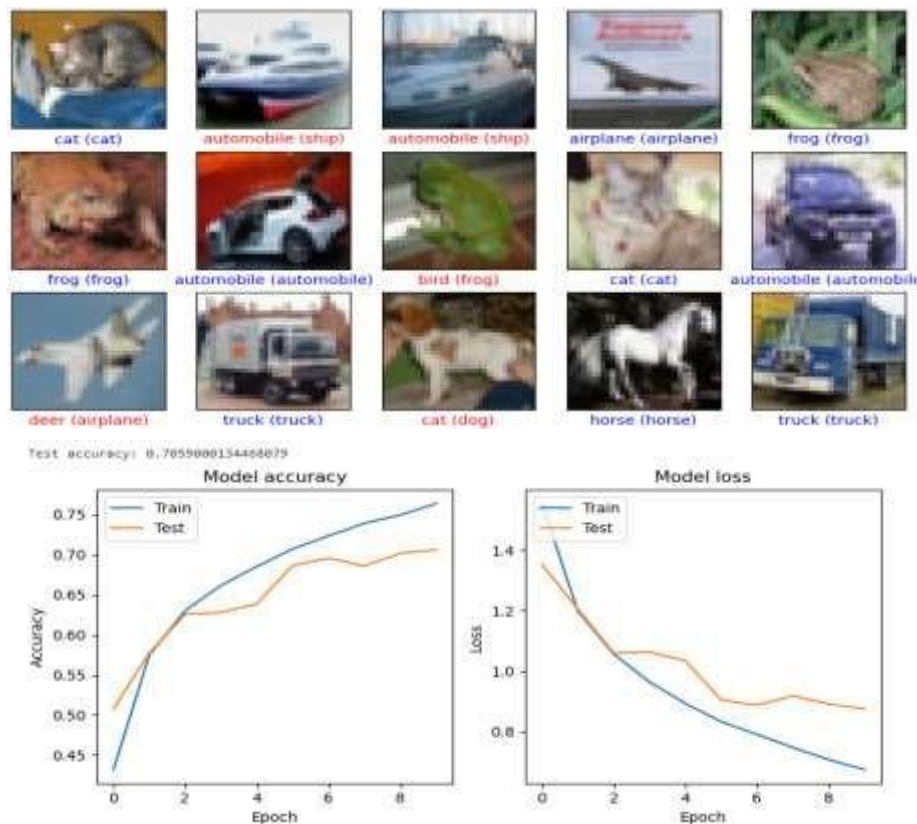
221501050

```

for i in range(20): plt.subplot(0, 0, i+1) plt.xticks([]) plt.yticks([]) plt.grid(False)
plt.imshow(X_test[i], cmap=plt.cm.binary) predicted_label = np.argmax(predictions[i])
true_label = y_test[i][0] color = 'blue' if predicted_label == true_label else 'red'
plt.xlabel(f'{class_names[predicted_label]} ({class_names[true_label]})',
color=color) plt.show()

```

### **OUTPUT:**



### **RESULT:**

Thus, the program for Image Classification is executed successfully.

## **7. PREDICTING DIABETES**

<b>EX.N0 : 7</b>	<b>PREDICTING DIABETES</b>
<b><u>DATE : 50/09/2024</u></b>	

**PROBLEM STATEMENT:** Predict the onset of diabetes based on medical measurements.

**PYTHON CONCEPTS:** Data structures, numeric types, functions.

**VISUALIZATION:** Scatter plots, heatmaps.

**MULTIVARIATE ANALYSIS:** Logistic regression, LDA.

**DATASET:** Pima Indians Diabetes Database **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

### **PROGRAM:**

```
import pandas as pd import seaborn as sns import matplotlib.pyplot as plt from
sklearn.model_selection import train_test_split from sklearn.linear_model import
LogisticRegression from sklearn.metrics import classification_report,
confusion_matrix, accuracy_score url =
https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indians-
diabetes.data.csv
columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
'DiabetesPedigreeFunction', 'Age', 'Outcome'] data =
pd.read_csv(url, header=None, names=columns) print("First 0
COMPUTATIONAL STATISTICS
```

221501050



```


records:\n", data.head()) print("\nStatistical Summary:\n",
data.describe()) print("\nDataset Info:\n") print(data.info())
sns.pairplot(data, hue='Outcome') plt.show() correlation_matrix
= data.corr() plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.show()
X = data.drop('Outcome', axis=1)
y = data['Outcome']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42) model
= LogisticRegression(max_iter=1500) model.fit(X_train, y_train) y_pred =
model.predict(X_test) print("Confusion Matrix:") print(confusion_matrix(y_test, y_pred))
print("\nClassification Report:") print(classification_report(y_test, y_pred)) accuracy =
accuracy_score(y_test, y_pred) print(f"\nModel Accuracy: {accuracy * 100:.2f}%") sample =
X_test.iloc[0].values.reshape(1, -1) sample_prediction = model.predict(sample)
print(f"\nPrediction for sample case (1 = Diabetes, 0 = No Diabetes): {sample_prediction[0]}")

```

### **OUTPUT:**

First 5 records:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	



Confusion Matrix:

```
[[120  31]
 [ 30  50]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.79	0.80	151
1	0.62	0.62	0.62	80
accuracy			0.74	231
macro avg	0.71	0.71	0.71	231
weighted avg	0.74	0.74	0.74	231

Model Accuracy: 73.59%

Prediction for sample case (1 = Diabetes, 0 = No Diabetes): 0

## RESULT:

Thus, the program for predicting diabetes is executed successfully.

## 8. WINE QUALITY PREDICTION

<b>EX.N0 : 8</b>	<b>WINE QUALITY PREDICTION</b>
<b><u>DATE : 18/09/2024</u></b>	

**PROBLEM STATEMENT:** Predict the quality of wine based on various chemical properties.

**PYTHON CONCEPTS:** Classes, sequences, file handling.

**VISUALIZATION:** Histograms, box plots.

**MULTIVARIATE ANALYSIS:** Multiple regression, factor analysis.

**DATASET:** Wine Quality Dataset **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

**PROGRAM:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
class WineQualityPredictor:
    def __init__(self, num_samples=1500):
```



```

self.num_samples = num_samples self.data = None self.model = None
def generate_data(self): np.random.seed(42) quality = np.random.randint(3, 9,
self.num_samples) # Quality scores between 3 and 8 fixed_acidity =
np.random.uniform(4.6, 10.9, self.num_samples) volatile_acidity =
np.random.uniform(0.12, 1.08, self.num_samples) citric_acid = np.random.uniform(0,
1, self.num_samples) residual_sugar = np.random.uniform(1.9, 10.0,
self.num_samples) chlorides = np.random.uniform(0.012, 0.1, self.num_samples)
free_sulfur_dioxide = np.random.uniform(1, 72, self.num_samples)
total_sulfur_dioxide = np.random.uniform(6, 289, self.num_samples) density =
np.random.uniform(0.99007, 1.00369, self.num_samples) pH =
np.random.uniform(2.74, 4.01, self.num_samples) sulfur_dioxide =
np.random.uniform(10, 60, self.num_samples) alcohol = np.random.uniform(8.0, 14.9,
self.num_samples) self.data = pd.DataFrame({
'fixed acidity': fixed_acidity, 'volatile acidity': volatile_acidity, 'citric acid': citric_acid,
'residual sugar': residual_sugar, 'chlorides': chlorides, 'free sulfur dioxide': free_sulfur_dioxide,
'total sulfur dioxide': total_sulfur_dioxide, 'density': density, 'pH': pH, 'sulphur
dioxide': sulfur_dioxide, 'alcohol': alcohol, 'quality': quality })
print(f'Synthetic Data Generated: {self.data.shape[0]} rows and {self.data.shape[1]} columns')
def visualize_data(self):
self.data.hist(bins=10, figsize=(10, 10))
plt.suptitle('Histograms of Wine Quality Features') plt.show()
plt.figure(figsize=(10, 6))
sns.boxplot(x='quality', y='fixed acidity',
data=self.data) plt.title('Box Plot of Fixed Acidity by
Quality') plt.show()
def preprocess_data(self): X =
self.data.drop('quality', axis=1) y = self.data['quality']
return X, y
def train_model(self, X, y):
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
self.model = LinearRegression() self.model.fit(X_train, y_train) y_pred =
self.model.predict(X_test) return y_train, y_test, y_pred
def evaluate_model(self,
y_test, y_pred): mse = mean_squared_error(y_test, y_pred) r2 = r2_score(y_test,
y_pred) print(f'Mean Squared Error: {mse}') print(f'R^2 Score: {r2}')
def predict_quality(self, input_features):
input_df = pd.DataFrame([input_features], columns=self.data.columns[:-
1]) prediction = self.model.predict(input_df) return prediction[0]
def

```

```
run(self): self.generate_data() self.visualize_data() X, y =  
self.preprocess_data() y_train, y_test, y_pred = self.train_model(X, y)  
self.evaluate_model(y_test, y_pred) if __name__ == "__main__":  
wine_predictor = WineQualityPredictor(num_samples=1500)  
wine_predictor.run() example_features = {  
'fixed acidity': 7.4, 'volatile acidity': 0.7, 'citric acid': 0.0,  
'residual sugar': 1.9, 'chlorides': 0.076, 'free sulfur dioxide': 50.0,  
'total sulfur dioxide': 34.0, 'density': 0.9978, 'pH': 3.01,  
'sulphur dioxide': 40.0, 'alcohol': 9.4 } predicted_quality =  
wine_predictor.predict_quality(example_features) print(f'Predicted  
Wine Quality: {predicted_quality:.2f}')
```

### **OUTPUT:**



Mean Squared Error: 2.8525212491984275  
 $R^2$  Score: -0.0010251435985495494  
 Predicted Wine Quality: 5.51

## **RESULT:**

Thus, the program for wine quality prediction is executed successfully.

## **9. HEART DISEASE PREDICTION**

<b>EX.N0 : 9</b>	<b>HEART DISEASE PREDICTION</b>
<b><u>DATE : 07/10/2024</u></b>	

**PROBLEM STATEMENT:** Predict heart disease based on clinical parameters **PYTHON**

**CONCEPTS:** Functions, data structures.

**VISUALIZATION:** Pair plots, ROC curve.

**MULTIVARIATE ANALYSIS:** Logistic regression, PCA.

**DATASET:** Heart Disease Dataset **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

**PROGRAM:**

```
import numpy as np import pandas as pd import matplotlib.pyplot as plt import
seaborn as sns from sklearn.model_selection import train_test_split from
sklearn.preprocessing import StandardScaler from sklearn.linear_model import
LogisticRegression from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report np.random.seed(42) # For reproducibility
num_samples = 1500 age = np.random.randint(30,
80, num_samples) sex = np.random.randint(0, 2,
num_samples) cp = np.random.randint(0, 4,
```

```

num_samples) trestbps = np.random.randint(90, 200,
num_samples) chol = np.random.randint(150, 300,
num_samples) fbs = np.random.randint(0, 2,
num_samples) restecg = np.random.randint(0, 2,
num_samples) thalach = np.random.randint(60, 200,
num_samples) exang = np.random.randint(0, 2,
num_samples) oldpeak = np.random.uniform(0, 6,
num_samples) slope = np.random.randint(0, 3,
num_samples) ca = np.random.randint(0, 4,
num_samples) thal = np.random.randint(1, 4,
num_samples) target = np.random.randint(0, 2,
num_samples) data = pd.DataFrame({
'age': age, 'sex': sex, 'cp': cp,
'trestbps': trestbps, 'chol': chol,
'fbs': fbs, 'restecg': restecg, 'thalach': thalach, 'exang': exang,
'oldpeak': oldpeak, 'slope': slope, 'ca': ca,
'thal': thal, 'target': target})
X = data.drop('target', axis=1) y
= data['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train) X_test =
scaler.transform(X_test) model =
LogisticRegression() model.fit(X_train, y_train)
y_pred = model.predict(X_test) accuracy =
accuracy_score(y_test, y_pred) conf_matrix =
confusion_matrix(y_test, y_pred) class_report =
classification_report(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}') print('Confusion
Matrix:') print(conf_matrix) print('Classification
Report:') print(class_report) plt.figure(figsize=(8,
6)) sns.heatmap(conf_matrix, annot=True, fmt='d',
cmap='Blues', xticklabels=['No Disease',
'Disease'], yticklabels=['No Disease', 'Disease']) plt.title('Confusion Matrix')
plt.xlabel('Predicted') plt.ylabel('Actual') plt.show() importance =

```



```
model.coef_[0] features = X.columns importance_df =  
pd.DataFrame({'Feature': features, 'Importance': importance}) importance_df =  
importance_df.sort_values(by='Importance', ascending=False)  
plt.figure(figsize=(10, 6)) sns.barplot(data=importance_df, x='Importance',  
y='Feature', palette='viridis') plt.title('Feature Importance')  
plt.xlabel('Coefficient Value') plt.ylabel('Features') plt.axvline(0, color='red',  
linestyle='--') # Adding a vertical line at 0 plt.show()
```

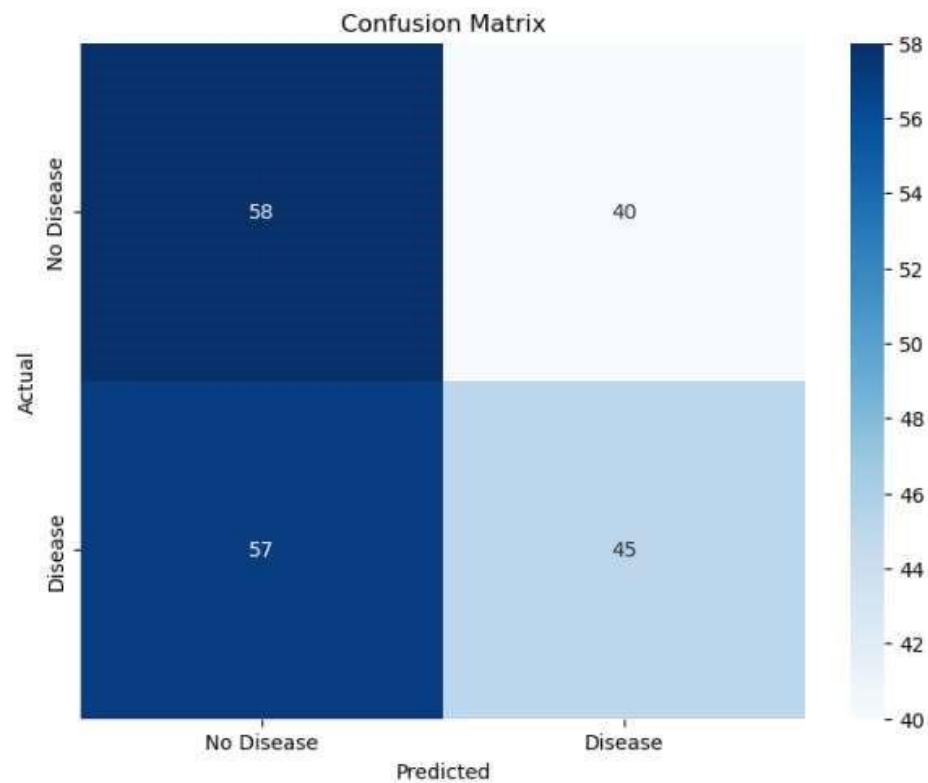
### **OUTPUT:**

```

Accuracy: 0.52
Confusion Matrix:
[[58 40]
 [57 45]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.50	0.59	0.54	98
1	0.53	0.44	0.48	102
accuracy			0.52	200
macro avg	0.52	0.52	0.51	200
weighted avg	0.52	0.52	0.51	200



## **RESULT:**

Thus, the program for heart disease prediction is executed successfully.

## **10. BREAST CANCER DIAGNOSIS**

<b>EX.N0 : 10</b>	<b>Breast Cancer Diagnosis</b>
-------------------	--------------------------------

**DATE : 09/10/2024**

**PROBLEM STATEMENT:** Classify tumors as benign or malignant based on features.

**PYTHON CONCEPTS:** Classes, sequences.

**VISUALIZATION:** Confusion matrix, bar plots.

**MULTIVARIATE ANALYSIS:** LDA, logistic regression.

**DATASET:** Breast Cancer Wisconsin Dataset **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

**PROGRAM:**

```
import numpy as np import pandas as pd import matplotlib.pyplot as plt import
seaborn as sns from sklearn.model_selection import train_test_split from
sklearn.preprocessing import StandardScaler from sklearn.linear_model import
LogisticRegression from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report np.random.seed(42) # For reproducibility
num_samples = 1500 age = np.random.randint(30,
80, num_samples) sex = np.random.randint(0, 2,
num_samples) cp = np.random.randint(0, 4,
num_samples) trestbps = np.random.randint(90, 200,
num_samples) chol = np.random.randint(150, 300,
num_samples) fbs = np.random.randint(0, 2,
```

```


num_samples) restecg = np.random.randint(0, 2,
num_samples) thalach = np.random.randint(60, 200,
num_samples) exang = np.random.randint(0, 2,
num_samples) oldpeak = np.random.uniform(0, 6,
num_samples) slope = np.random.randint(0, 3,
num_samples) ca = np.random.randint(0, 4,
num_samples) thal = np.random.randint(1, 4,
num_samples) target = np.random.randint(0, 2,
num_samples) data = pd.DataFrame({
'age': age, 'sex': sex, 'cp': cp,
'trestbps': trestbps, 'chol': chol,
'fbs': fbs, 'restecg': restecg, 'thalach': thalach, 'exang': exang,
'oldpeak': oldpeak, 'slope': slope, 'ca': ca,
'thal': thal, 'target': target})
X = data.drop('target', axis=1) y
= data['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train) X_test =
scaler.transform(X_test) model =
LogisticRegression() model.fit(X_train, y_train)
y_pred = model.predict(X_test) accuracy =
accuracy_score(y_test, y_pred) conf_matrix =
confusion_matrix(y_test, y_pred) class_report =
classification_report(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}') print('Confusion
Matrix:') print(conf_matrix) print('Classification
Report:') print(class_report) plt.figure(figsize=(8,
6)) sns.heatmap(conf_matrix, annot=True, fmt='d',
cmap='Blues', xticklabels=['No Disease',
'Disease'], yticklabels=['No Disease', 'Disease']) plt.title('Confusion Matrix')
plt.xlabel('Predicted') plt.ylabel('Actual') plt.show() importance =
model.coef_[0] features = X.columns importance_df =
pd.DataFrame({'Feature': features, 'Importance': importance}) importance_df =
importance_df.sort_values(by='Importance', ascending=False)

```

```
plt.figure(figsize=(10, 6)) sns.barplot(data=importance_df, x='Importance',  
y='Feature', palette='viridis') plt.title('Feature Importance')  
plt.xlabel('Coefficient Value') plt.ylabel('Features') plt.axvline(0, color='red',  
linestyle='--') # Adding a vertical line at 0 plt.show()
```

### **OUTPUT:**

	precision	recall	f1-score	support
0	0.97	0.99	0.98	71
1	0.98	0.95	0.96	43
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114



```

Enter the following features for prediction: compactness_se: 0.03
radius_mean: 14.5 concavity_se: 0.03
texture_mean: 20.0 concave points_se: 0.02
perimeter_mean: 90.0 symmetry_se: 0.02
area_mean: 560.0 fractal_dimension_se: 0.003
smoothness_mean: 0.1 radius_worst: 16.0
compactness_mean: 0.15 texture_worst: 25.0
concavity_mean: 0.2 perimeter_worst: 100.0
concave points_mean: 0.1 area_worst: 800.0
symmetry_mean: 0.18 smoothness_worst: 0.14
fractal_dimension_mean: 0.06 compactness_worst: 0.25
radius_se: 0.6 concavity_worst: 0.3
texture_se: 1.2 concave points_worst: 0.15
perimeter_se: 10.0 symmetry_worst: 0.25
area_se: 40.0 fractal_dimension_worst: 0.08
smoothness_se: 0.007
The tumor is predicted to be: Malignant
Based on the symptoms provided, the person may be at risk.

```

## **RESULT:**

Thus, the program for breast cancer diagnosis is executed successfully.

## **50. PREDICTING FLIGHT DELAYS**

<b>EX.N0 : 50</b>	<b>PREDICTING FLIGHT DELAYS</b>
<b><u>DATE : 16/10/2024</u></b>	

**PROBLEM STATEMENT:** Predict flight delays based on historical data.

**PYTHON CONCEPTS:** File reading/writing, functions.

**VISUALIZATION:** Line plots, scatter plots.

**MULTIVARIATE ANALYSIS:** Regression, clustering.

**DATASET:** Flight Delay Dataset **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

**PROGRAM:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

df = pd.read_csv('C:/Users/APPU/Downloads/Airline_Delay_Cause.csv')
print(df.columns)
```

```

print(df.isnull().sum()) df.dropna(inplace=True) # or
df.fillna(method='ffill', inplace=True) if 'year' in df.columns and
'month' in df.columns:
df['date'] = pd.to_datetime(df[['year', 'month']].assign(day=1))
plt.figure(figsize=(10, 0)) sns.lineplot(data=df, x='date',
y='arr_delay') # Adjust if necessary plt.title('Flight Delays Over
Time') plt.xticks(rotation=40) plt.show() delay_column =
'arr_delay' # Using 'arr_delay' for now if 'carrier_delay' in
df.columns and delay_column in df.columns:
plt.figure(figsize=(10, 0)) sns.scatterplot(data=df, x='carrier_delay',
y=delay_column) # Adjust as needed plt.title('Carrier Delay vs Arrival Delays')
plt.xlabel('Carrier Delay (minutes)') plt.ylabel('Arrival Delay (minutes)')
plt.show()
else: print("Check the delay columns: 'carrier_delay' or 'arr_delay' do not exist in the
DataFrame.") df['day_of_week'] = df['date'].dt.dayofweek # Monday=0, Sunday=6
features = ['day_of_week', 'arr_flights', 'carrier_ct'] # Modify as needed
X = df[features] y = df[delay_column]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) model
= LinearRegression() model.fit(X_train, y_train) predictions = model.predict(X_test)
print('Mean Absolute Error:', mean_absolute_error(y_test, predictions)) print('Mean Squared
Error:', mean_squared_error(y_test, predictions)) print('R-squared:', r2_score(y_test,
predictions)) plt.figure(figsize=(10, 0)) plt.scatter(y_test, predictions) plt.plot([min(y_test),
max(y_test)], [min(y_test), max(y_test)], color='red', linewidth=2) # Line of equality
plt.title('Predictions vs Actual Delays') plt.xlabel('Actual Delays') plt.ylabel('Predicted Delays')
plt.show()



```



## OUTPUT:

```
Index(['year', 'month', 'carrier', 'carrier_name', 'airport', 'airport_name',  
      'arr_flights', 'arr_del15', 'carrier_ct', 'weather_ct', 'nas_ct',  
      'security_ct', 'late_aircraft_ct', 'arr_cancelled', 'arr_diverted',  
      'arr_delay', 'carrier_delay', 'weather_delay', 'nas_delay',  
      'security_delay', 'late_aircraft_delay'],  
      dtype='object')
```

```
year          0  
month         0  
carrier       0  
carrier_name  0  
airport       0  
airport_name  0  
arr_flights   240  
arr_del15     443  
carrier_ct    240  
weather_ct    240  
nas_ct        240  
security_ct   240  
late_aircraft_ct 240  
arr_cancelled 240  
arr_diverted  240  
arr_delay     240  
carrier_delay 240  
weather_delay 240  
nas_delay     240  
security_delay 240  
late_aircraft_delay 240  
dtype: int64
```



Mean Absolute Error: 1592.2201262853362  
Mean Squared Error: 25524907.35571326  
R-squared: 0.8439698040165798

COMPUTATIONAL STATISTICS

221501050

## RESULT:

## **12. ENERGY CONSUMPTION FORECASTING**

<b>EX.N0 : 12</b>	<b>ENERGY CONSUMPTION FORECASTING</b>
<b><u>DATE : 23/10/2024</u></b>	

**PROBLEM STATEMENT:** Forecast energy consumption based on historical data.

**PYTHON CONCEPTS:** Functions, numeric types.

**VISUALIZATION:** Line plots, heatmaps.

**MULTIVARIATE ANALYSIS:** Time series analysis, regression.

**DATASET:** Energy Consumption Dataset **ALGORITHM:**

Step 1: Start the program.

Step 2: Import necessary libraries.

Step 3: Load the dataset.

Step 4: Encode categorical variable, define feature & testing set.

Step 0: Split the dataset into training & testing set, create trained model.

Step 6: Print equal metric & test the cell.

### **PROGRAM:**



```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
data = pd.read_csv('C:/Users/APPU/Downloads/energy_consumption_dataset.csv', parse_dates=['Timestamp'],
index_col='Timestamp')
print(data.head())
```

---

```
print(data.info()) data = data.fillna(method='ffill')
plt.figure(figsize=(14, 6))
plt.plot(data['EnergyConsumption'], color='blue',
label='Energy Consumption') plt.title('Energy
Consumption Over Time') plt.xlabel('Date')
plt.ylabel('Consumption') plt.legend() plt.show()
numeric_data =
data.select_dtypes(include=[np.number])
plt.figure(figsize=(10, 8))
sns.heatmap(numeric_data.corr(), annot=True,
cmap='coolwarm') plt.title('Correlation Matrix')
plt.show() from statsmodels.tsa.seasonal import
seasonal_decompose result =
seasonal_decompose(data['EnergyConsumption'],
model='additive', period=24) # Adjust period based on
your data's frequency result.plot() plt.show() train_size
= int(len(data) * 0.8) train, test =
data['EnergyConsumption'][:train_size],
data['EnergyConsumption'][train_size:] model =
ARIMA(train, order=(0, 1, 0)) # Adjust (p,d,q) based
on your data's behavior fitted_model = model.fit()
forecast = fitted_model.forecast(steps=len(test))
forecast_index = test.index mse =
mean_squared_error(test, forecast) rmse = np.sqrt(mse)
print(f'RMSE: {rmse}') plt.figure(figsize=(14, 6))
plt.plot(train, label='Train') plt.plot(test, label='Test')
plt.plot(forecast_index, forecast, label='Forecast')
plt.title('Energy Consumption Forecast')
plt.xlabel('Date') plt.ylabel('Consumption') plt.legend()
plt.show()
```

## OUTPUT:

Timestamp	Temperature	Humidity	SquareFootage	Occupancy	Timestamp	HVACUsage	LightingUsage	RenewableEnergy	DayOfWeek
2022-01-01 00:00:00	25.139433	43.431581	1565.693999	5	2022-01-01 00:00:00	On	Off	2.774699	Monday
2022-01-01 01:00:00	27.731651	54.225919	1411.064918	1	2022-01-01 01:00:00	On	On	21.831384	Saturday
2022-01-01 02:00:00	28.704277	58.907658	1755.715009	2	2022-01-01 02:00:00	Off	Off	6.764672	Sunday
2022-01-01 03:00:00	20.000469	50.371637	1452.316318	1	2022-01-01 03:00:00	Off	On	8.623447	Wednesday
2022-01-01 04:00:00	23.097359	51.401421	1094.130359	9	2022-01-01 04:00:00	On	Off	3.071969	Friday



COMPUTATIONAL STATISTICS

221501050

## RESULT:

Thus, the program for energy consumption forecasting is executed successfully.